

Why Deep Learning rocks

A deep philosophical note

Maxim Borisyak

Yandex Machine Intelligence and Research
National Research University Higher School of Economics

No free lunch

IQ test: try to learn yourself!

First question from MENSA website:

Following the pattern shown in the number sequence below, what is the missing number?

1, 8, 27, ?, 125, 216

Possible answers:

- › 36
- › 45
- › 46
- › 64
- › 99

IQ test: try to learn yourself!

First question from MENSA website:

Following the pattern shown in the number sequence below, what is the missing number?

X_{train}	1	2	3	5	6
y_{train}	1	8	27	125	216

$$X_{\text{test}} = (4,)$$

IQ test: try to learn yourself!

My solution:

$$y = \frac{1}{12}(91x^5 - 1519x^4 + 9449x^3 - 26705x^2 + 33588x - 14940)$$

› fits perfectly!

My answer:

› 99

Terminology

Machine Learning is about learning algorithms A that:

- › defined on sample set \mathcal{X} (e.g. \mathbb{R}^n) and targets \mathcal{Y} (e.g. $\{0, 1\}$);
- › take a problem (dataset) $D = (X, y) \subseteq \mathcal{X} \times \mathcal{Y}$;
- › learn relation between \mathcal{X} and \mathcal{Y} ;
- › and return prediction function:

$$\begin{aligned} A(D) &= f \\ f : \mathcal{X} &\rightarrow \mathcal{Y} \end{aligned}$$

By this definition, e.g. XGBoost is a **family** of algorithms.

No free lunch theorem

No free lunch theorem states that **on average by all datasets** all learning algorithms are equally bad at learning.


Examples:

› crazy algorithm:




$$f(x) = \left[\left(\left[\sum_i x_i + \theta \right] \bmod 17 + 1027 \right)^\pi \right] \bmod 2$$

› any configuration of SVM
perform equally well **on average**.

No free lunch theorem, stat. edition



No free lunch theorem



Possible learning algorithm behaviours in **problem space**:

- › $+$ - better than the average;
- › $-$ - worse than the average.

Are Machine Learning algorithms useless?

Are Machine Learning algorithms useless?

No.

Are Machine Learning algorithms useless?

- › No Free Lunch theorem applies to:
 - › one learning algorithm;
 - › against all possible problems.
- › in real world we have:
 - › **data scientist** with prior knowledge of the world;
 - › problem description;
 - › data description;
 - › a set of standard algorithms.

Corollary

A good machine learning family of algorithms/framework:


- › has clear relation between hyperparameters and set of problems each algorithm covers;
- › i.e. a data scientist can easily map their prior knowledge on hyperparameters.

A great machine learning family/frameworks:

- › covers a wide range of problems;
- › but each algorithm covers a small set of problems;
- › i.e. a lot of sensitive and well-defined hyperparameters.

Traditional Machine Learning (simplified)

- › analyse the problem and make assumptions;
- › pick an algorithm from a toolkit (e.g. logistic regression);
- › provide assumptions suitable for the algorithm (**feature engineering**).



Discussion

- › this approach works well for traditional datasets with a small number of features:
- › e.g. Titanic dataset:

passenger class	name	gender	age	fare	...
-----------------	------	--------	-----	------	-----

Essentially, performance of the algorithm depends on:

- › knowledge of the domain;
- › feature engineering skills;
- › understanding of assumptions behind standard algorithms.

Kitten

Let's try to detect kittens!




Kitten seen by a machine

```
[[ 22  25  28  32  29 ..., 58  36  35  34  34]
 [ 26  29  30  31  36 ..., 65  38  42  41  42]
 [ 27  28  31  30  40 ..., 84  58  51  52  44]
 [ 27  26  27  29  43 ..., 90  70  60  57  43]
 [ 20  26  28  28  31 ..., 83  73  62  52  45]
 ...,
 [173 187 180 183 184 ..., 170 227 244 219 199]
 [193 199 194 188 185 ..., 181 197 201 209 187]
 [175 177 156 166 171 ..., 226 215 194 185 182]
 [161 159 160 187 178 ..., 216 193 220 211 200]
 [178 180 177 185 164 ..., 190 184 212 216 189]]
```

Solution?

- › edge detection;
- › image segmentation;
- › eyes, ears, nose models;
- › fit nose, ears, eyes;
- › average color of segments;
- › standard deviation of color segments;
- › goodness of fit for segments;
- › kitten's face model;
- › logistic regression.

Solution?



Solution?



Perhaps, more Machine Learning and less Human Engineering?

Deep Learning



Deep Learning

Let's learn features!

Deep Learning



Deep Learning



Kitten

Traditional approach:

- › edge detection;
- › image segmentation;
- › fit nose, ears, eyes;
- › average, standard deviation of segment color;
- › fluffiness model;
- › kitten's face model;
- › logistic regression.

Deep Learning:

- › non-linear transformation;
- › another non-linear transformation;
- › non-linear transformation, again;
- › non-linear transformation, and again;
- › non-linear transformation (why not?);
- › logistic regression.

Deep Learning

- › is not a superior algorithm;
- › is not even a single algorithm;
- › is a framework;
- › allows to express our assumptions in much more general way.

Why DL rocks

- › can crack much harder problems;
 - › it is easier to formulate models for features than features itself;
- › easy to construct networks:
 - › merge together;
 - › bring new objectives;
 - › inject something inside network;
 - › build networks inside networks;
 - › any differentiable magic is allowed*.

* Non-differentiable also, but with a special care.

Example

A problem contains groups of features:


- › image;
- › sound features;

Prior knowledge:

- › features from different group should not interact directly;

Example of a solution:

- › build a subnetwork upon each group of features;
- › merge them together.



Almost Free Lunch

Disclaimer

This is not a comprehensive overview of Deep Learning, just some examples.

Hacking layers


- › restrictions on weights: convolutions, weight matrix decomposition, ...;
- › activation: ReLU, ELU, SELU, ...;
- › new operations: pooling, maxout, ...;
- › specific unit behaviour: GRU, LSTM units, ...

Hacking model

Restrictions on search space:


› regularization, e.g.:

$$\mathcal{L} = \mathcal{L}_{\text{cross-entropy}} + \alpha \|W\|_2^2$$



Hacking search procedure

- › SGD-like methods:
 - › adam, adadelata, adamax, rmsprop;
 - › nesterov momentum;
- › quasi-Newton methods;
- › batch normalization, weight normalization;
- › weight matrix decomposition.




Data augmentation

- › symmetries: shifts, rotations, ...:
 - › searching for a network that produces the same response for shifted/rotated samples;
 - › eliminating symmetries;
- › random noise:
 - › pushing separation surface farther from samples - robust output;


Hacking model

Interference with network (change of objective):

- › drop-out, drop-connect:
- › searching for a robust network.



(a) Standard Neural Net



(b) After applying dropout.

Hacking loss

Hacking objectives:

- › introducing loss for each layer:

$$\mathcal{L} = \mathcal{L}_n + \sum_{i=1}^{n-1} C_i \mathcal{L}_i$$


where:

- › \mathcal{L}_i - loss on i -th layer.
- › Deeply Supervised Networks:
 - › searches for network that obtains good intermediate results.

Hacking initial guess


Pretraining:

- › unsupervised pretraining;
- › greedy layer-wise pretraining with e.g. AutoEncoders.



Hacking architecture: U-net

- › skip connections allow to combine context with low-level representation.



Hacking architecture: ResNet

- › residual connections produce boosting-like behaviour;
- › no vanishing gradients;
- › brute-force: up to 1000+ layers.







Figure 2. Residual learning: a building block.


Hacking architecture: Cycle-GAN

- › reverse transformation;
- › symmetric discrimination.



Hacking architecture: Ladder Network

- › introduces auxiliary task: denoising;



Breaking No Free Lunch

Auxiliary task

- › introducing additional task for a network might considerably improve generalization:

$$\mathcal{L} = \mathcal{L}_{\text{main}} + \alpha \mathcal{L}_{\text{auxiliary}}$$

- › e.g. along with particle identification restore momentum;
- › brings additional information.

Auxiliary
losses

Main
loss



Pretraining

- › pretraining on a similar dataset;
 - › as initial guess;
 - › regularization relative to another model $\{W_i^0\}_{i=1}^n$:

$$\mathcal{L} = \mathcal{L}_{\text{main}} + \lambda \sum_{i=1}^n \|W_i - W_i^0\|_2^2$$

- › using weights from the zoo;
- › unsupervised pretraining (semi-supervised learning).

Summary

Summary

No Free Lunch theorem:

- › Machine Learning is about using prior knowledge about the problem wisely.

Deep Learning:

- › a framework that covers a large range of problems;
- › allows to express prior knowledge freely;
- › makes it easier to solve hard problems.

References

No-Free-Lunch theorem:

- › Schaffer, Cullen. "A conservation law for generalization performance." Proceedings of the 11th international conference on machine learning. 1994.
- › Wolpert, David H. "The supervised learning no-free-lunch theorems." Soft computing and industry. Springer London, 2002. 25-42.
- › Wolpert, David H., and William G. Macready. "No free lunch theorems for optimization." IEEE transactions on evolutionary computation 1.1 (1997): 67-82.

References: layers

- › Clevert, Djork-Arné, Thomas Unterthiner, and Sepp Hochreiter. "Fast and accurate deep network learning by exponential linear units (elus)." arXiv preprint arXiv:1511.07289 (2015).
- › Goodfellow, Ian J., et al. "Maxout networks." arXiv preprint arXiv:1302.4389 (2013).
- › Chung, Junyoung, et al. "Gated feedback recurrent neural networks." International Conference on Machine Learning. 2015.

References: regularization

- › Sainath, Tara N., et al. "Low-rank matrix factorization for deep neural network training with high-dimensional output targets." Acoustics, Speech and Signal Processing (ICASSP), 2013 IEEE International Conference on. IEEE, 2013.
- › Srivastava, Nitish, et al. "Dropout: a simple way to prevent neural networks from overfitting." Journal of Machine Learning Research 15.1 (2014): 1929-1958.
- › Lee, Chen-Yu, et al. "Deeply-supervised nets." Artificial Intelligence and Statistics. 2015.

References: network architectures

- › Von Eicken, Thorsten, et al. "U-Net: A user-level network interface for parallel and distributed computing." ACM SIGOPS Operating Systems Review. Vol. 29. No. 5. ACM, 1995.
- › He, Kaiming, et al. "Deep residual learning for image recognition." Proceedings of the IEEE conference on computer vision and pattern recognition. 2016.
- › Zhu, Jun-Yan, et al. "Unpaired image-to-image translation using cycle-consistent adversarial networks." arXiv preprint arXiv:1703.10593 (2017).
- › Rasmus, Antti, et al. "Semi-supervised learning with ladder networks." Advances in Neural Information Processing Systems. 2015.

More resources

A lot of useful links can be found in:

- › Schmidhuber, Jürgen. "Deep learning in neural networks: An overview." Neural networks 61 (2015): 85-117.
- › Bengio, Yoshua. "Learning deep architectures for AI." Foundations and trends® in Machine Learning 2.1 (2009): 1-127.